

Assignment 1

Fabio Zanotti, Antonio Morelli, Federico Rullo, and Edoardo Conca

Master's Degree in Artificial Intelligence, University of Bologna

{ fabio.zanotti, antonio.morelli, federico.rullo, edoardo.conca }@studio.unibo.it

Abstract

In this project, we present the implementation and evaluation of a Part-Of-Speech (POS) tagging system using neural architectures. We explore three different neural models: a baseline model with a single Bidirectional Long Short-Term Memory (LSTM) layer, an extended model with an additional Bidirectional LSTM layer, and a variant with an extra Dense layer. The results showed how the deeper models behaved better than the baseline, although the simpler architecture performs comparably to the more complex models, which may suggest that the complexity introduced by the extra layers did not provide substantial benefits for the task.

1 Introduction

The assignment aims to develop a POS tagging system using neural architectures to predict the POS tag for each word in a corpus. Three models are implemented and compared to determine performance superiority. Bidirectional LSTM layers are utilized for contextual understanding, crucial for accurate POS tagging. Methodology involves model implementation and comparison using the average Macro-F1 score as a performance metric. Preprocessing includes dataset encoding and integration of GloVe embeddings for weight initialization, leveraging pre-trained representations. Out-Of-Vocabulary (OOV) words are incorporated into the vocabulary, with embeddings computed based on tag-specific means. Experimental evaluation indicates improved accuracy in models with additional Bidirectional LSTM or Dense layers, emphasizing enhanced contextual and feature representation in POS tagging.

2 System description

The three different models implemented are:

- Baseline - consisting of a single Bidirectional LSTM layer;

- Baseline with an additional Bidirectional LSTM layer, to further capture contextual dependencies;
- Baseline with an additional Dense layer, to enhance the feature representation;

We started by importing the dataset from the Natural Language Toolkit(NLTK) library, which we pre-processed and then encoded, and removing the "-NONE-" tag, so that the model will have fewer examples of unstructured data to learn, and can focus on the examples that are most relevant to the POS tagging, increasing the accuracy of the results. After inspecting the dataset composition, which seemed to be affected by some imbalance in between classes, we built the vocabulary; for this, we used the 100-dimensional version of GloVe to balance computational efficiency and semantic richness. Subsequently, these embeddings initialized the model's weights, leveraging pre-trained word representations to enhance our task's performance. To address Out-Of-Vocabulary (OOV) words, the GloVe vocabulary was expanded by integrating OOV words from the dataset. Instead of just assigning random vectors to OOV words, tag-specific means from existing words were computed, ensuring meaningful embeddings and improving model performance. This process was iteratively applied to the training, validation, and test sets to compute the vocabulary.

Encoded sentences were padded to the length of the second-longest sentence, excluding the longest one because it turned out to be an outlier, and the padding label was defined as "-PAD-", following treebank notation.

Finally, we implemented the three models for the POS Tagging system. Each of these models have the same input and output layers, comprising GloVe Embedding as the input layer and TimeDistributed as the output layer. The differentiating factor among these models lies in their hidden layers. The

initial model, designated as the Baseline model, consists solely of a single Bidirectional LSTM hidden layer 1. The second model extends the Baseline model by incorporating an additional Bidirectional LSTM layer 2. And lastly, the third and ultimate model integrates an additional Dense layer into the Baseline architecture 3.

3 Experimental setup and results

3.1 Parameter Tuning

Prior to conducting experiments, hyperparameter tuning was performed for each model. We observed significant effects on scores due to variations in LSTM layer units and batch size. Balancing model complexity and overfitting, we set the LSTM units between 128 and 256. A batch size of 32 was chosen to enhance performance, reduce variance, and address class imbalance. Consistency across models was ensured by employing the same parameters. Callbacks were utilized for effective training control.

3.2 Metrics

Custom metrics were defined to accommodate non-informative classes in POS tagging evaluation. Classes such as punctuation and symbols were ignored, with custom metrics including a masked F1 macro score for post-training evaluation. To address dataset imbalance, a custom accuracy metric was devised. This metric assigns different weights to errors on specific classes and computes accuracy accordingly. By excluding certain classes and adjusting weights, the metric provides a balanced evaluation of model performance.

3.3 Results

The following tables provide a clear comparison of model performances on both validation and test sets, highlighting the effectiveness of the Additional Dense model in improving macro F1 scores.

Table 1: Avg F1 Scores on Validation Set

	Macro F1 Scores
Baseline Model	0.7891
Additional LSTM	0.7966
Additional Dense	0.7992

Table 2: Avg F1 Scores on Test Set

	Macro F1 Scores
Baseline Model	0.7810
Additional LSTM	0.7836
Additional Dense	0.7851

4 Discussion

4.1 Quantitative Results

Upon analyzing the results, it is evident that while the performances of the two deeper models are comparable, the one with two Dense layers exhibits slightly superior performance on average. However, adhering to the outlined request for error analysis on the best single model based on the validation set performance, we focused on one of the LSTM models, which achieved a macro F1-score of 0.8054. This observation suggests that LSTM layers effectively captured sequential dependencies and long-range relationships within the input data, crucial for POS tagging tasks.

4.2 Error analysis

We observed a higher error rate on the validation set compared to the test set, possibly due to more challenging or edge case samples in the validation set. Four classes—NN, NNP, JJ, and NNS—were consistently among the top mismatches for each model, with their higher presence in the validation set explaining the increased error rate. Investigating further, we found that errors predominantly occurred between classes with semantic and syntactic similarities, such as NN (Noun, singular or mass), NNP (Proper noun, singular), JJ (Adjective), and NNS (Noun, plural). Confusion between these classes stems from their overlapping characteristics, such as singular vs. plural forms and common vs. proper nouns.

5 Conclusion

Surprisingly, we found that the simpler Baseline model performances are worse by a marginal factor of approximately $\approx 1\%$. This suggests that the additional layers in the deeper models did not significantly improve performance. One possible explanation for this is that the complexity introduced by the extra layers did not provide substantial benefits for the POS tagging task. It's possible that the Baseline model's simpler architecture was already effective in capturing the necessary patterns in the data, leading to comparable performance despite

its simplicity

6 Links to external resources

[Assignment_1](#)

Appendix

Figures

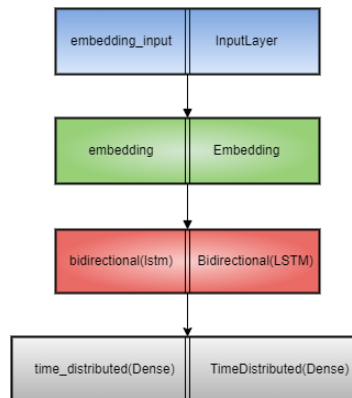


Figure 1: Baseline architecture

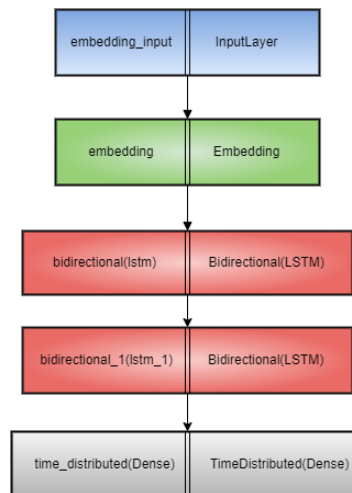


Figure 2: Additional LSTM architecture

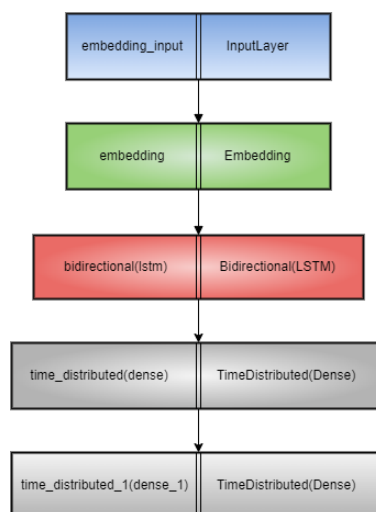


Figure 3: Additional Dense architecture